Citation Provenance Verification Paper

This is the introductory page for citation provenance verification testing. It introduces speech recognition models and scaling benchmarks.

Figure 8. Zero-shot Whisper performance scales reliably across tasks and languages with increasing model size.

Table 6. Performance improves with increasing dataset size.

All increases in the dataset size result in improved performance on all tasks, although we see significant variability in improvement rates across tasks and sizes. Performance improves rapidly on English speech recognition from 3,000 to 13,000 hours and then slows down noticeably between 13,000 and 54,000 hours.

Multitask learning provides robust transfer across speech recognition and translation benchmarks with minimal fine-tuning required.

We evaluate models with -1 and +1 score adjustments across 1-3 training phases.

Identical repeated passage on page two for ambiguity testing. Identical repeated passage on page two for ambiguity testing.